

User Manual: Harmonic Tide Prediction

Script name: `tide.py` (*Python 3*)

Abstract

This manual explains how to run, and interpret the outputs of a non-professional harmonic prediction tool for sea level that generates a short animation. The program reads an hourly time series in CSV format (`year,month,day,hour,level_mm`), converts levels from millimeters to meters, splits the record into two halves (training and test), fits a tidal harmonic model by least squares on the first half, and, on the second half, animates sliding 7-day windows comparing observed vs. predicted levels. It also computes and plots 4-week moving averages over the full period and saves an animation as MP4 via `ffmpeg`.

This software is not for professional use. It was created using `vibe` programming to illustrate how harmonic prediction can be done using standard libraries in Python.

Contents

1 Overview	2
2 Requirements	2
3 Input data (CSV)	2
3.1 Format	2
3.2 File location	2
4 Quick start	3
5 Command-line options	3
6 Harmonic model and fitting	4
6.1 Constituent set	4
6.2 Frequencies	4
6.3 Design matrix and least squares	4
7 Animation: sliding 7-day window	4
7.1 Windowing and frames	4
7.2 Visual style	5
7.3 Video output	5
8 Four-week moving averages	5
9 Logging	5
10 Key validations and messages	5
11 Customization and extensions	5
11.1 Change constituents	5
11.2 Window length	5
11.3 Styling	5
11.4 GIF export	6
12 Best practices and reproducibility	6
13 Troubleshooting (FAQ)	6

1 Overview

The software implements the harmonic model

$$y(t) = \beta_0 + \sum_{k=1}^K (A_k \cos(\omega_k t) + B_k \sin(\omega_k t)),$$

where $y(t)$ is sea level (in meters) relative to a mean offset, ω_k are angular frequencies in rad/s derived from standard speeds in degrees per hour, and (β_0, A_k, B_k) are estimated by least squares. The frequency set follows the NOS/CO-OPS *standard suite* (about 37 constituents including M2, S2, N2, K1, O1, etc.), together with shallow-water and long-period terms as encoded in SPEEDS_DPH.

This software is not for professional use. It was created using `vibe` programming to illustrate how harmonic prediction can be done using standard libraries in Python.

Do not use for navigation or operational decisions.

2 Requirements

- **Python 3.8+**
- Python packages: `numpy`, `pandas`, `matplotlib`
- `ffmpeg` installed and on your PATH (for saving the animation)

3 Input data (CSV)

3.1 Format

The CSV must have **five columns without a header**:

`year, month, day, hour, level_mm`

- `year, month, day, hour` are integers in UTC.
- `level_mm` is an integer in millimeters; the program converts it to meters.
- Missing values are flagged by `-32767 mm` (i.e., `-32.767 m`) and are **removed**.

The data available in the folder `./data/` for the cities Honolulu, Fortaleza and Salvador were obtained for University of Hawai'i Sea Level Center (UHSLC) <https://uhslc.soest.hawaii.edu/data/> and the license described there for this data applies. If you use UHSLC tide gauge data in your research or applications, please cite the dataset as:

Caldwell, P. C., M. A. Merrifield, P. R. Thompson (2015), Sea level measured by tide gauges from global oceans — the Joint Archive for Sea Level holdings (NCEI Accession 0019568), Version 5.5, NOAA National Centers for Environmental Information, Dataset, doi:10.7289/V5V40S7W.

3.2 File location

- If `-csv_file` is **not** provided, the program lists `.csv` files under `./data/` and prompts you to choose one.
- If `-csv_file` is provided, the path must exist; otherwise execution aborts with an error.

4 Quick start

Minimal example 1 (Interactive):

```
1 python tide.py
```

Minimal example 2:

```
1 python tide.py --csv_file ./data/Honolulu.csv --initial_year 1920 --
  final_year 2000 \
2 --week_seed 42 --animation_weeks 4 --local "Honolulu" --
  animation_speed 1.0
```

On launch, the program:

1. Reads and sorts the time series; reports available year range.
2. If `-initial_year`/`-final_year` are missing, prompts for years interactively.
3. Removes missing values (-32767) and checks that at least 100 samples remain.
4. Splits data into two halves: training (first half) and test (second half).
5. Subtracts the *offset* (training mean) from both halves.
6. Fits harmonic coefficients by least squares on the training set.
7. Predicts over the test period and starts a 7-day sliding-window animation.
8. Computes & plots 4-week moving averages and saves a PNG.
9. Saves an MP4 animation using `ffmpeg`.

5 Command-line options

Option	Type / Default	Description
<code>-csv_file</code>	str / (none)	Path to the CSV file. If omitted, the program scans <code>./data/</code> and asks interactively.
<code>-initial_year</code>	int / (available minimum)	Start year of the time slice. Validated not to be before the earliest year. If both years are omitted, selection is interactive.
<code>-final_year</code>	int / (available maximum)	End year of the time slice. Validated not to exceed the latest year and to be \geq start year.
<code>-week_seed</code>	int / 42	RNG seed for the initial window position in the test half. Ensures reproducible starting week.
<code>-animation_weeks</code>	int / 4	Animation duration in <i>weeks</i> (limits the number of frames). Increase to cover more of the test period.
<code>-local</code>	str / ""	Station/location name shown in titles and used to compose the video filename.
<code>-animation_speed</code>	float / 1.0	Speed multiplier for the animation (affects the per-frame step in samples). Larger values advance faster. PS: use 1.0 (there is a bug here)

Practical examples.

1. Interactive CSV and year selection:

```
1 python tide.py
```

2. Years 1998–2015, 12-week animation, 1x speed:

```
1 python tide.py --csv_file ./data/Honolulu.csv --initial_year 1998 --  
  final_year 2015 \  
2 --animation_weeks 12 --animation_speed 1.0 --local "Honolulu"
```

3. Reproduce the same initial week with a fixed seed:

```
1 python tide.py --csv_file ./data/Fortaleza-Brazil.csv --week_seed 7
```

6 Harmonic model and fitting

6.1 Constituent set

The script defines a dictionary of speeds in degrees per hour (SPEEDS_DPH) with ≈ 37 standard constituents (M2, S2, N2, K2, K1, O1, P1, Q1, ...), shallow-water harmonics (M4, MS4, MN4, M6, M8), long-period terms (SA, SSA, MF, MM, MSF), and some additional components. You may edit this list directly to tailor the model.

6.2 Frequencies

Each speed v (deg/h) is converted to a period T in hours via $T = 360/v$. The angular frequency in rad/s is

$$\omega = \frac{2\pi}{T \cdot 3600}.$$

Time is referenced to the first training timestamp t_0 using $t = (\text{timestamp} - t_0)$ in seconds.

6.3 Design matrix and least squares

The design matrix contains a column of ones and, for each ω_k , the columns $\cos(\omega_k t)$ and $\sin(\omega_k t)$. Parameters θ are estimated by `numpy` least squares:

$$\theta = \arg \min_{\theta} \|X\theta - y\|_2^2.$$

The training mean (offset) is subtracted from both training and test so that predictions and observations share the same baseline in plots.

7 Animation: sliding 7-day window

7.1 Windowing and frames

- The *test* half is traversed with 7-day windows.
- The number of samples per window is derived from the data timestep Δt .
- The initial position is drawn using `-week_seed`.
- `-animation_weeks` limits the total number of frames.
- `-animation_speed` scales the per-frame shift (in samples).

7.2 Visual style

The script uses a dark theme (`plt.style.use('dark_background')`), solid lines for observations and dashed lines for predictions, a subtle grid, and titles summarizing the total data period, current window, and number of constituents. The y -axis limits adapt to each window with a 10% margin.

7.3 Video output

The animation is built with `matplotlib.animation.FuncAnimation` and saved via `ffmpeg`. Default filename rules:

- If `-local` is provided: `{local_lower}_weekly_average.mp4`.
- Otherwise: `{csv_basename}_weekly_average_{YYYYMMDD_HHMM}.mp4`.

8 Four-week moving averages

After the animation, the program computes 4-week averages (with a 1-week stride for time markers) across the *entire* record, using the same offset removed during training. It also fits a linear trend in time (years) and displays the slope in m/year in the legend. The figure is saved as:

`{csv_basename}_average_level_4_weeks_{YYYYMMDD_HHMM}.png`.

9 Logging

Events are recorded to `tide.log`: full command line, chosen CSV, filtered period, record counts, number of 4-week averages, basic statistics (mean, std), and, if applicable, the trend slope.

10 Key validations and messages

- **Missing CSV**: execution stops with an error.
- **Invalid year range**: `start < available minimum`, `end > available maximum`, or `start > end` → error.
- **After missing-value filtering**: if fewer than 100 records remain, execution stops with an error.

11 Customization and extensions

11.1 Change constituents

Edit `SPEEDS_DPH` to add/remove constituents. Closely spaced frequencies can cause ill-conditioning; consider regularization or selecting a subset.

11.2 Window length

To animate windows other than 7 days, change `win_days` in the code and recompute `win_samples` accordingly.

11.3 Styling

Colors, line widths, and fonts can be adjusted where `line_r`, `line_p`, axes, and legends are defined.

11.4 GIF export

Switching the writer to `PillowWriter` allows saving GIFs, though MP4 via `ffmpeg` usually yields better quality and size.

12 Best practices and reproducibility

- Fix `-week_seed` to reproduce the same starting window.
- Record the exact command line (also logged in `tide.log`).

13 Troubleshooting (FAQ)

Error: File '...' not found.

Check the `-csv_file` path or place the CSV under `./data/`.

No CSV files found during interactive run

Create `data/` and move your CSV(s) there.

Insufficient data after filtering

After removing missing values, fewer than 100 rows remain.

Video not saved / ffmpeg error

Ensure `ffmpeg` is installed and on your PATH.

Unreadable colors/labels

Adjust the theme colors or increase `fontsize` in the plotting section.

Quick reference

Listing 1: Common commands

```
1 # 1) Full run with labels and 8-week animation
2 python tide.py --csv_file ./data/Honolulu.csv --initial_year 1993 --
   final_year 2020 \
3   --local "Fortaleza-Brazil" --animation_weeks 8 --animation_speed 1.0
4
5 # 2) Interactive run (no args): select CSV and years at the prompt
6 python tide.py
7
8 # 3) Faster animation and fixed starting week
9 python tide.py --csv_file data/Fortaleza-Brazil.csv --week_seed 123 --
   animation_speed 1.0
```

Credits and license

The license for the code is MIT.

The data available in the folder `./data/` for the cites Honolulu, Fortaleza and Salvador were obtained for University of Hawai'i Sea Level Center (UHSLC) <https://uhslc.soest.hawaii.edu/data/> and the license described there for this data applies. If you use UHSLC tide gauge data in your research or applications, please cite the dataset as:

Caldwell, P. C., M. A. Merrifield, P. R. Thompson (2015), Sea level measured by tide gauges from global oceans — the Joint Archive for Sea Level holdings (NCEI Accession 0019568), Version 5.5, NOAA National Centers for Environmental Information, Dataset, doi:10.7289/V5V40S7W.

This manual was last updated on October 4, 2025.